Spectral Subsampling MCMC for Stationary Multivariate Time Series

Mattias Villani

Department of Statistics
Stockholm University

mattiasvillani.com @matvil @matvil mattiasvillani

Overview

- **Background: Subsampling MCMC/HMC**
- **Spectral subsampling for multivariate time series**
- **Application to vector ARTFIMA**

- **Slides**: http://mattiasvillani.com/news.

Joint work with some Aussies

- Subsampling MCMC for time series:
  - **Robert Kohn**, UNSW Sydney
  - **Matias Quiroz**, UTS Sydney
  - **Robert Salomone**, QUT Brisbane
- Subsampling MCMC/HMC (conditionally) independent data:
  - **Minh-Ngoc Tran**, Univ of Sydney
  - **Khue-Dung Dang**, Univ of Melbourne
- Subsampling for spatial data:
  - **Tom Goodwin**, UNSW Sydney
  - **Arthur P. Guillaumin**, Queen Mary University of London

This talk: spectral subsampling for time series

- Salomone, Quiroz, Kohn, Villani and Tran (2020)
  Spectral Subsampling MCMC for Stationary Time Series
  International Conference on Machine Learning (*ICML*).

- Villani, Quiroz, Kohn and Salomone (2024)
  Spectral Subsampling MCMC for Stationary Multivariate Time Series with Applications to Vector ARTFIMA Processes
  Econometrics & Statistics, *Part B Statistics*.

Motivation

- **Long time series** are increasingly **common**:
  - high frequency financial transaction data
  - neuroimaging data with high temporal resolution
  - sensor data from robots
  - meteorological weather stations
  - GPS data used in urban traffic monitoring.
- Often **multivariate observations**.
- Automatic **decision making under uncertainty**.
- **Bayesian decisions**: maximize **posterior expected utility**.
- Posteriors by **MCMC/HMC** simulation.
- **MCMC/HMC is slow** on large datasets. 😔

The Metropolis-Hastings (MH) algorithm

- **Bayesian inference**

$$p(\boldsymbol{\theta}|\mathbf{y}) \propto L(\boldsymbol{\theta})p(\boldsymbol{\theta})$$

- Initialize $\boldsymbol{\theta}^{(0)}$ and iterate for $k = 1, 2, ..., N$

  1. Sample $\boldsymbol{\theta}_p \sim q\left(\cdot|\boldsymbol{\theta}^{(k-1)}\right)$ (the **proposal distribution**)
  2. Accept $\boldsymbol{\theta}_p$ with **acceptance probability**

$$\alpha = \min\left(1, \frac{L(\boldsymbol{\theta}_p)p(\boldsymbol{\theta}_p)}{L\left(\boldsymbol{\theta}^{(k-1)}\right)p\left(\boldsymbol{\theta}^{(k-1)}\right)} \frac{q\left(\boldsymbol{\theta}^{(k-1)}|\boldsymbol{\theta}_p\right)}{q\left(\boldsymbol{\theta}_p|\boldsymbol{\theta}^{(k-1)}\right)}\right)$$

- **Costly** to evaluate $L(\boldsymbol{\theta}_p)$ when n is large. **Big data**.

Naive Subsampling MH

- Independent data - **log-likelihood** is a **sum**

$$\ell(\boldsymbol{\theta}) \stackrel{\text{def}}{=} \log L(\boldsymbol{\theta}) = \sum_{i=1}^{n} \log \ell_i(\boldsymbol{\theta}), \qquad \ell_i(\boldsymbol{\theta}) \stackrel{\text{def}}{=} p(y_i|\boldsymbol{\theta})$$

- **Unbiased estimate** of $\ell(\boldsymbol{\theta})$ from **subsample** of size $m \ll n$

$$\hat{\ell}(\boldsymbol{\theta}, \mathbf{u}) = \frac{n}{m} \sum_{i \in \mathbf{u}} \log \ell_i(\boldsymbol{\theta})$$

- Run **Pseudo-marginal MH** with $\hat{L}(\boldsymbol{\theta}, \mathbf{u}) = \exp(\hat{\ell}(\boldsymbol{\theta}, \mathbf{u}))$.

- Initialize $(\boldsymbol{\theta}^{(0)}, \mathbf{u}^{(0)})$ and iterate for $k = 1, 2, ..., N$
  1. Sample $\boldsymbol{\theta}_p \sim q(\cdot|\boldsymbol{\theta}^{(k-1)})$ and subsample $\mathbf{u}_p \sim p(\mathbf{u})$
  2. Accept $(\boldsymbol{\theta}_p, \mathbf{u}_p)$ with **acceptance probability**

$$\alpha = \min\left(1, \frac{\hat{L}(\boldsymbol{\theta}_p, \mathbf{u}_p) p(\boldsymbol{\theta}_p)}{\hat{L}(\boldsymbol{\theta}^{(k-1)}, \mathbf{u}^{(i-1)}) p(\boldsymbol{\theta}^{(k-1)})} \frac{q(\boldsymbol{\theta}^{(k-1)}|\boldsymbol{\theta}_p)}{q(\boldsymbol{\theta}_p|\boldsymbol{\theta}^{(k-1)})}\right)$$

Fixing Naive Subsampling MH - Bias

- **If $\hat{L}$ unbiased** then samples are from $p(\boldsymbol{\theta}|\boldsymbol{y})$ [1]
- Unbiased log-likelihood: $\mathbb{E}_{\mathbf{u}}[\hat{\ell}(\boldsymbol{\theta},\mathbf{u})] = \ell(\boldsymbol{\theta})$
- But biased likelihood estimate: $\hat{L}(\boldsymbol{\theta},\mathbf{u}) = \exp(\hat{\ell}(\boldsymbol{\theta},\mathbf{u}))$
- **Approximate bias correction of** $\exp\left(\hat{\ell}(\boldsymbol{\theta},\mathbf{u})\right)$ [2]

  Theorem: Error in posterior approximation is $O\left(\frac{1}{m^2 n}\right)$ [3]
- **Unbiased Block-Poisson estimator** + **Signed PMMH** [4]

Fixing Naive Subsampling MH - Variance

- **Low** $\mathbb{V}\left(\hat{L}\left(\boldsymbol{\theta}, \mathbf{u}\right)\right)$ is crucial for **efficient sampling**.
- **Difference estimator** with control variates [3]

$$\widehat{\ell}_{\mathrm{diff}}(\boldsymbol{\theta}, \mathbf{u}) := \sum_{k=1}^{n} q_k(\boldsymbol{\theta}) + \frac{n}{m} \sum_{i=1}^{m} \left(\ell_{u_i}(\boldsymbol{\theta}) - q_{u_i}\left(\boldsymbol{\theta}\right)\right)$$

- **Control variates** $q_{u_i}\left(\boldsymbol{\theta}\right)$ by Taylor expansion around $\tilde{\boldsymbol{\theta}}$. [3, 5]
- **Optimal tuning** of subsample size m [6, 3, 4]
- **Blocking**: only refresh part of the subsample [7, 8]
- **Grouping observations** for improved control variates [9]
- **High-dim θ**: **Subsampling HMC**. [10]

Beyond independent data

- Subsampling methods **assume** the **log-likelihood is a sum**

$$\ell(\boldsymbol{\theta}) = \sum_{i=1}^{n} \log p(y_i|\boldsymbol{\theta})$$

- Estimating $\ell(\boldsymbol{\theta})$ is like estimating a **population total**

$$\hat{\ell}(\boldsymbol{\theta}, \boldsymbol{u}) = \frac{n}{m} \sum_{i \in \boldsymbol{u}} \log p(y_i|\boldsymbol{\theta})$$

- **Log-likelihood is a sum:**
  - for conditionally independent y_i
  - for longitudinal data when subjects are independent.
  - for special time series, e.g. AR processes. Sample (x_t, x_{t-1}).
- General **time series** dependence? Spatial dependence?

Spectral density of a stationary process

- **Autocovariance function**

$$\gamma(\tau) = \mathbb{E}\left[(x_t - \mu)(x_{t-\tau} - \mu)\right], \quad \tau = 0, 1, \ldots$$

- **Spectral density**

$$f(\omega) \equiv \frac{1}{2\pi} \sum_{\tau=-\infty}^{\infty} \gamma(\tau) \exp(-\mathrm{i}\omega\tau) \text{ for } \omega \in (-\pi, \pi].$$

- **AR(1) process:** $x_t = \phi x_{t-1} + \varepsilon_t, \quad \varepsilon_t \overset{iid}{\sim} N(0, \sigma_\varepsilon^2)$

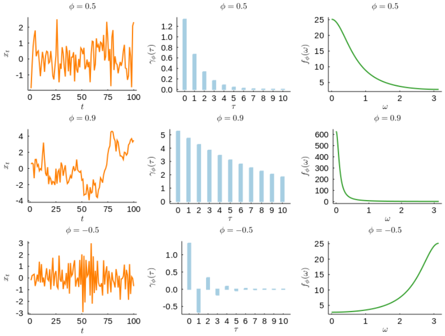

Discrete Fourier transform

- **Discrete Fourier Transform** (DFT) of the time series

$$J(\omega_k) \equiv \frac{1}{\sqrt{2\pi}} \sum_{t=1}^{n} x_t \exp(-\mathrm{i}\omega_k t)$$

- The **periodogram**

$$\mathcal{I}(\omega_k) = n^{-1} \left| J(\omega_k) \right|^2 .$$

- **Asympotically** as $n \to \infty$

$$\mathcal{I}(\omega_k) \overset{\text{indep}}{\sim} \text{Exponential}\left(f\left(\omega_k\right)\right), \quad k = 1, \ldots, n$$

- **Whittle's** asymptotic approximation of the log-likelihood:

$$\ell_W(\boldsymbol{\theta}) \equiv -\sum_{\omega_k \in \Omega} \left(\log f_{\boldsymbol{\theta}}(\omega_k) + \frac{\mathcal{I}(\omega_k)}{f_{\boldsymbol{\theta}}(\omega_k)} \right)$$

- **Whittle log-likelihood is a sum**. **Subsample frequencies!**

Multivariate Fourier analysis

- **Autocovariance matrix function** for time series $\boldsymbol{x}_t \in \mathbb{R}^r$

$$\gamma_{\boldsymbol{x}}(\tau) = \mathrm{Cov}(\boldsymbol{x}_t, \boldsymbol{x}_{t-\tau}) = \left[\gamma_{jk}(\tau)\right]_{i,j=1,\ldots,r}$$

- **Spectral density matrix**

$$f_{\boldsymbol{x}}(\omega) = \frac{1}{2\pi}\sum_{\tau=-\infty}^{\infty} \gamma_{\boldsymbol{x}}(\tau)\exp(-i\omega\tau)$$

where off-diagonal elements are the **cross-spectral densities**

$$f_{jk}(\omega) = \sum_{\tau=-\infty}^{\infty} \gamma_{jk}(\tau)\exp(-i\omega\tau)$$

- **Multivariate** discrete Fourier transform (**DFT**)

$$J(\omega_k) = \sum_{t=0}^{n-1} \boldsymbol{x}_t \exp(-i\omega_k t)$$

Subsampling MCMC for multivariate time series

- DFT asymptotically independent complex multiv normal [11]

$$n^{-1/2}J(\omega_k) \overset{\text{indep}}{\sim} \mathrm{CN}(0, 2\pi f_{\boldsymbol{x}}(\omega_k)) \text{ as } n \to \infty.$$

- **Multivariate periodogram** is complex singular Wishart

$$I_T(\omega) = (2\pi n)^{-1} J(\omega) J_T(\omega)^H \sim \mathrm{CW}(1, f_{\boldsymbol{x}}(\omega))$$

- **Multivariate Whittle** log-likelihood [12]

$$\ell_{\mathcal{W}}(\boldsymbol{\theta}) = -\sum_{\omega_k \in \Omega_n} \left(\log |f_{\boldsymbol{x}}(\omega_k)| + \mathrm{tr}\left[f_{\boldsymbol{x}}(\omega_k)^{-1} I_T(\omega)\right]\right)$$

- Whittle **biased for small n**
- ... but **subsampling** only relevant for **large n**.

Univariate ARTFIMA

- **ARFIMA**(p, d, q) with fractional differencing d

$$\phi_p(L)(1-L)^{d} x_t = \theta_q(L)\varepsilon_t$$

- **Long memory**. $\sum_{\tau=-\infty}^{\infty} |\gamma(\tau)| = \infty$. But stationary if $|d| < 1/2$.

$$(1-L)^d x_t \stackrel{\text{def}}{=} \sum_{j=0}^{\infty} (-1)^j \frac{\Gamma(1+d)}{\Gamma(1+d-j)j!} x_{t-j}$$

- **ARTFIMA** adds **tempering** parameter $\lambda \geq 0$ [13]

$$(1-e^{-\lambda}L)^d x_t \stackrel{\text{def}}{=} \sum_{j=0}^{\infty} (-1)^j \frac{\Gamma(1+d)}{\Gamma(1+d-j)j!} e^{-\lambda j} x_{t-j}$$

  - long range dependence in $\gamma(\tau)$ for small τ
  - exponential decay for larger τ
  - Stationary for all d and $\lambda > 0$.

Vector ARTFIMA(p, d, λ, q)

- Multivariate extension of ARTFIMA for r-dim $\boldsymbol{x}_t$ [12]

$$\Phi_p(L)\Delta^{\boldsymbol{d},\boldsymbol{\lambda}}(\boldsymbol{x}_t - \boldsymbol{\mu}) = \Theta_q(L)\boldsymbol{\varepsilon}_t, \quad \boldsymbol{\varepsilon}_t \overset{\text{iid}}{\sim} N(0, \Sigma_\varepsilon)$$

where

$$\Delta^{\boldsymbol{d},\boldsymbol{\lambda}} \equiv \mathrm{Diag}\big((1 - e^{-\lambda_1}L)^{d_1}, \ldots, (1 - e^{-\lambda_r}L)^{d_r}\big)$$

- VARTFIMA is **stationary** and causal for all $\boldsymbol{d}$ and $\boldsymbol{\lambda} > \mathbf{0}$.
- **Spectral density matrix**

$$f_{\boldsymbol{x}}(\omega) = \frac{1}{2\pi}\boldsymbol{B}\Phi_p^{-1}(e^{-i\omega})\Theta_q(e^{-i\omega})\Sigma_\varepsilon\Theta_q(e^{-i\omega})\Phi_p^{-H}(e^{-i\omega})^H\boldsymbol{B}^H$$

$$\boldsymbol{B} = \mathrm{Diag}\big((1 - e^{-(\lambda_1 + i\omega)})^{-d_1}, \ldots, (1 - e^{-(\lambda_r + i\omega)})^{-d_r}\big).$$

- **Ansley-Kohn parametrization** of both Φ and Θ to ensure stationarity and invertibility.
- Aim: joint posterior

$$p(\Phi, \Theta, \boldsymbol{d}, \boldsymbol{\lambda} | \boldsymbol{x}_{1:n})$$

Three datasets for evaluation

- **Swedish temperatures**
  - Three locations: Arlanda, Bromma and Landvetter.
  - Hourly data from February 1, 2008 until May 1, 2022.
- **Water velocity**
  - Mean water velocity every 12th minute at two locations on opposite sides of Lake St Clair.
  - 130,001 observations from Jan 3, 2016 until Dec 21, 2018.
- **Air pollution** in Stockholm
  - Nitrogen dioxide (NO2) and particulate matter (PM10) pollution at two streets in central Stockholm.
  - Hourly data for the time period February 16, 2010 until October 31, 2015.
- **Subsample**: **1%** of sample, using control variates for groups.

Water velocity data

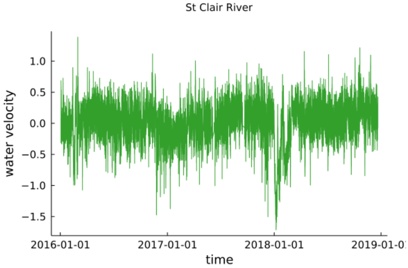

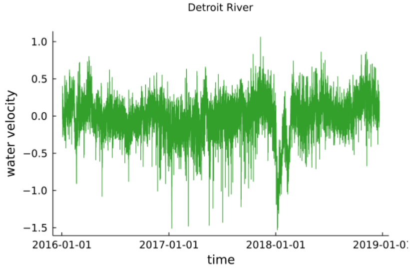

Model selection via BIC approximation

		Water Velocity		Temperature		Pollution	
AR	MA	No TFI	TFI	No TFI	TFI	No TFI	TFI
1	0	737079	759123	327097	334122	363760	366022
0	1	588297	759457	61320	332888	306068	365658
2	0	749650	761200	335201	**335757**	365522	**366266**
0	2	621765	**761786**	93256	333948	325717	366142
1	1	758838	761305	333582	335647	365762	**366267**

Computional times

- **Computational Time (CT)**

$$\text{CT} = \text{Inefficiency factor} \times \text{Compute time for single draw}$$

- **Relative Computational Time (RCT)**:

$$\text{RCT} = \frac{\text{CT MCMC full data sample}}{\text{CT Spectral subsampling MCMC}}$$

Dataset	Model	Min	Mean	Max
Water velocity	VARTFIMA(0,2)	87	98	125
Temperature	VARTFIMA(2,0)	68	89	114

Variance of log-likehood estimator is crucial

- Spectral subsampling can fail when $Var(\hat{\ell})$ is too large.
- VARTFIMA(2,0) for Swedish temperature data:

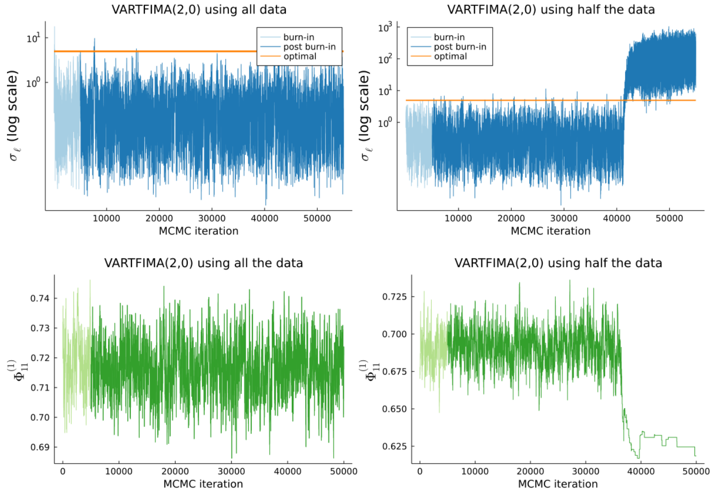

VARTFIMA(0,2) - Subsampling is accurate

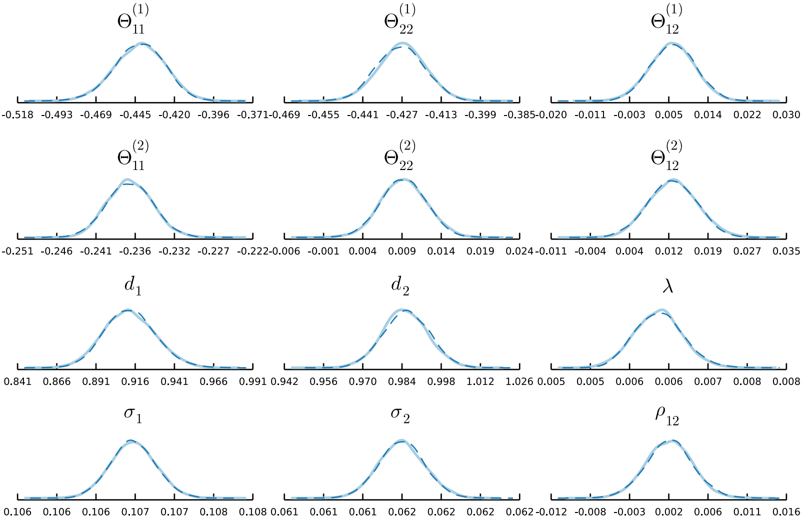

VARMA(1,1) - Whittle is accurate

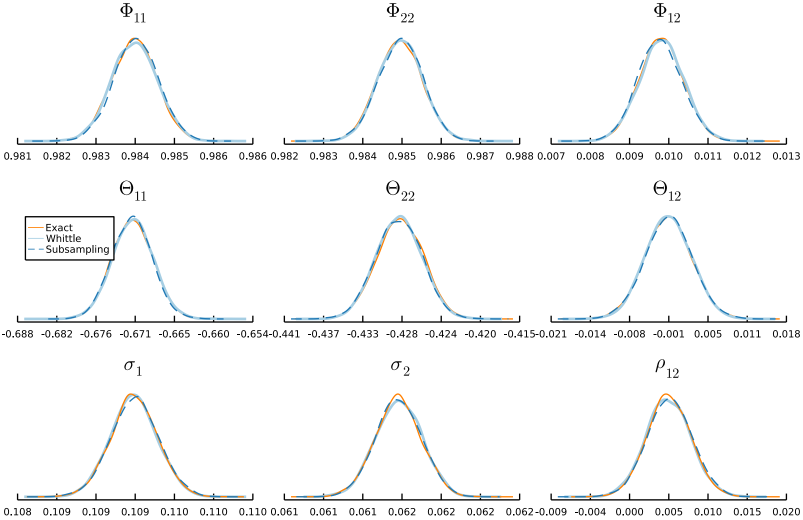

VARTFIMA(0,2) - good model fit

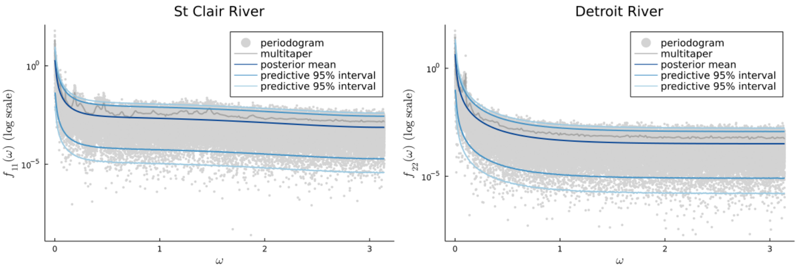

VARTFIMA(0,2) - coherence and delay (phase)

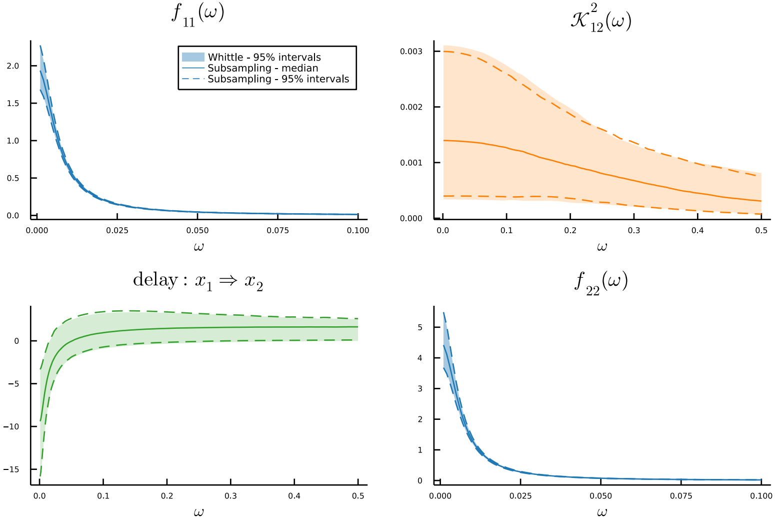

Swedish temperature data

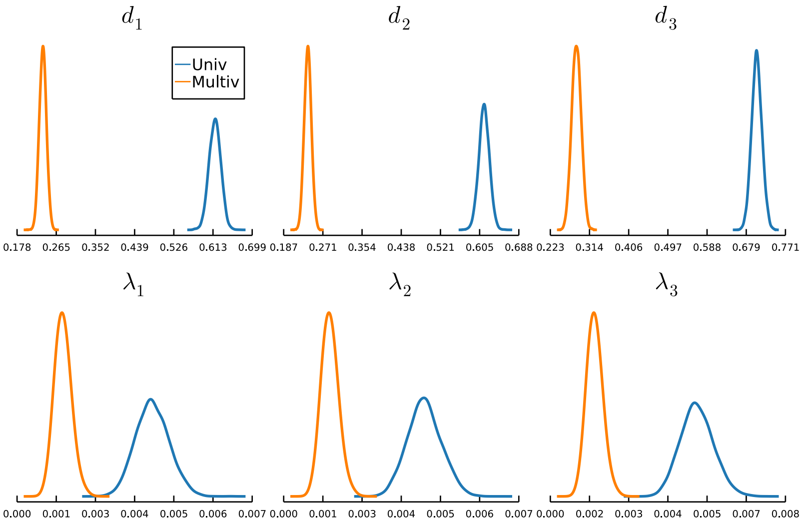

Conclusions

- **Whittle log-likelihood** is fast to compute and is a sum.
- **Whittle enables subsampling** for time series.
- **Subsampling of matrix periodogram** data to speed up MCMC/HMC for multivariate time series.
- **Very large speed-ups** compared to regular MCMC/HMC.
- In progress: **Spatial data** with **debiased Whittle**.

References

- C. Andrieu and G. O. Roberts, "The pseudo-marginal approach for efficient Monte Carlo computations," *The Annals of Statistics*, pp. 697–725, 2009.
- D. Ceperley and M. Dewing, "The penalty method for random walks with uncertain energies," *The Journal of chemical physics*, vol. 110, no. 20, pp. 9812–9820, 1999.
- M. Quiroz, R. Kohn, M. Villani, and M.-N. Tran, "Speeding up mcmc by efficient data subsampling," *Journal of the American Statistical Association*, no. 114, 2019.
- M. Quiroz, M.-N. Tran, M. Villani, R. Kohn, and K.-D. Dang, "The block-Poisson estimator for optimally tuned exact subsampling MCMC," *Journal of Computational and Graphical Statistics*, 2021.
- R. Bardenet, A. Doucet, and C. Holmes, "On markov chain monte carlo methods for tall data," *The Journal of Machine Learning Research*, vol. 18, no. 1, pp. 1515–1557, 2017.

- M. K. Pitt, R. d. S. Silva, P. Giordani, and R. Kohn, "On some properties of Markov chain Monte Carlo simulation methods based on the particle filter," *Journal of Econometrics*, vol. 171, no. 2, pp. 134–151, 2012.
- M.-N. Tran, R. Kohn, M. Quiroz, and M. Villani, "Block-wise pseudo-marginal metropolis-hastings," *arXiv preprint arXiv:1603.02485*, 2016.
- G. Deligiannidis, A. Doucet, and M. K. Pitt, "The correlated pseudo-marginal method," *arXiv preprint arXiv:1511.04992*, 2015.
- R. Salomone, M. Quiroz, R. Kohn, M. Villani, and M.-N. Tran, "Spectral subsampling mcmc for stationary time series," *ICML2020*, 2020.
- K.-D. Dang, M. Quiroz, R. Kohn, M.-N. Tran, and M. Villani, "Hamiltonian monte carlo with energy conserving subsampling," *Journal of Machine Learning Research, 2019*, vol. 20, pp. 1–31, 2019.

- D. R. Brillinger, *Time series: data analysis and theory*. SIAM, 2001.
- M. Villani, M. Quiroz, R. Kohn, and R. Salomone, "Spectral subsampling mcmc for stationary multivariate time series with applications to vector artfima processes," *Econometrics and Statistics*, 2024.
- F. Sabzikar, A. I. McLeod, and M. M. Meerschaert, "Parameter estimation for ARTFIMA time series," *Journal of Statistical Planning and Inference*, vol. 200, pp. 129 – 145, 2019.